

Here's a simplified end-to-end explanation of the project:

---

### ### **\*\*Overview\*\***

This project is a **\*\*PDF Question Answering System\*\*** that lets users upload PDF files, extract text from them, and ask natural language questions about their content. The system processes the PDFs, identifies the most relevant text, and generates a response using an AI model like Claude.

### ### **\*\*Key Features\*\***

1. **\*\*PDF Upload\*\***: Users can upload multiple PDF files through a user-friendly Streamlit interface.
2. **\*\*Text Extraction\*\***: The system extracts text content from the uploaded PDFs.
3. **\*\*Embedding Creation\*\***: Text is split into smaller chunks and converted into numerical representations (embeddings) using a Hugging Face model.
4. **\*\*Vector Store\*\***: These embeddings are stored in a vector database (FAISS) for efficient similarity search.
5. **\*\*Question Answering\*\***: Users input a question, and the system retrieves the most relevant text chunks, formats a prompt, and sends it to the Claude AI model for generating answers.
6. **\*\*Interactive UI\*\***: Streamlit provides an intuitive interface for all these functionalities.

---

### ### **\*\*Step-by-Step Flow\*\***

#### #### 1. **\*\*Initialization\*\***

- **\*\*Environment Setup\*\***: Loads API keys and other configurations from environment variables.
- **\*\*Dependencies\*\***:
  - ``PyPDF2``: Extracts text from PDF files.
  - ``langchain``: Manages text splitting and embeddings.
  - ``Hugging Face Embeddings``: Converts text into embeddings for semantic search.

- `FAISS`: Handles vector-based similarity search.
- `Streamlit`: Provides a web-based user interface.
- `Anthropic`: Communicates with the Claude AI model.

#### #### 2. **User Interaction**

- The user uploads one or more PDF files through the Streamlit interface.
- Files are temporarily stored in a local directory for processing.

#### #### 3. **Text Processing**

- **Text Extraction**: Text is extracted from all uploaded PDFs using `PyPDF2`.
- **Splitting**: Long text is split into smaller chunks (e.g., 1000 characters with overlaps) for better search accuracy.
- **Embedding Generation**: Each text chunk is transformed into a numerical format using a Hugging Face embedding model.
- **Vector Store Creation**: The embeddings are saved in a FAISS vector store for efficient retrieval.

#### #### 4. **Question Answering**

- **Query Input**: The user types a question in the Streamlit input box.
- **Similarity Search**: The system searches the FAISS vector store for the most relevant text chunks.
- **Prompt Creation**: The relevant chunks are combined into a prompt formatted for the Claude AI model.
- **Response Generation**: The prompt is sent to the Claude API, which generates an answer based on the context provided.

#### #### 5. **Output**

- The AI-generated answer is displayed in the Streamlit app.
- If the system cannot find enough context to answer, it informs the user accordingly.

---

### ### **\*\*Technology Stack\*\***

#### 1. **\*\*Backend\*\***:

- `Python`
- `Claude AI API` (Anthropic)
- `FAISS` for vector search
- `Hugging Face` for embeddings

#### 2. **\*\*Frontend\*\***:

- `Streamlit` for the web interface

#### 3. **\*\*PDF Handling\*\***:

- `PyPDF2` for text extraction

---

### ### **\*\*How to Use\*\***

1. **\*\*Upload PDFs\*\***: Use the sidebar to upload PDF files.
2. **\*\*Process PDFs\*\***: Click the "Process PDFs" button to extract and prepare the data.
3. **\*\*Ask Questions\*\***: Enter a question in the text box. The system will fetch the most relevant chunks and provide an answer.

---

### ### **\*\*Business Value\*\***

- **\*\*Efficiency\*\***: Quickly find answers from lengthy documents without manually reading them.
- **\*\*Scalability\*\***: Handles multiple PDFs and large datasets.
- **\*\*Interactivity\*\***: Easy-to-use web interface accessible to non-technical users.

---

### ### **\*\*Example Use Case\*\***

1. **\*\*Upload\*\***: Add PDFs containing company reports or policy documents.
2. **\*\*Process\*\***: Extract and organize the document contents.
3. **\*\*Query\*\***: Ask questions like:
  - "What is the main conclusion in these reports?"
  - "Summarize the risk factors discussed in the documents."

---

Let me know if you'd like clarification on any part or further enhancements!

\*\*\*\*\*  
\*\*\*\*\*

Sure! Let's go through the code step by step, explaining each line in simple terms:

### **\*\*1. Importing Libraries:\*\***

```
```python
import os
import faiss
import pickle

from anthropic import Anthropic
from PyPDF2 import PdfReader
from langchain.text_splitter import RecursiveCharacterTextSplitter
from langchain_community.embeddings import HuggingFaceEmbeddings
from langchain_community.vectorstores import FAISS
from dotenv import load_dotenv

import streamlit as st

import textwrap
```

...

- **os**: Used for interacting with the operating system, like reading files or creating directories.
- **faiss**: A library to efficiently search through large datasets, like finding similar text chunks.
- **pickle**: Used to save and load Python objects (like our FAISS index) between runs of the program.
- **Anthropic**: Used to interact with the Claude AI model for generating answers.
- **PdfReader**: A tool to read PDF files and extract text.
- **RecursiveCharacterTextSplitter**: Helps to split large chunks of text into smaller parts, making it easier to process.
- **HuggingFaceEmbeddings**: A tool for converting text into numerical embeddings (representations) that machine learning models can use.
- **FAISS**: A tool that allows you to store and search for similar text chunks.
- **load\_dotenv**: Loads environment variables (like API keys) from a `.env` file.
- **streamlit**: A framework used to create interactive web apps.
- **textwrap**: Used to format the text output, making it easier to read.

### **2. Loading Environment Variables:**

```
python
```

```
load_dotenv()
```

...

- This loads environment variables (like API keys) from a `.env` file into the program. This is important for keeping sensitive information (like API keys) secure.

### **3. The PDFQuestionAnswerer Class:**

```
python
```

```
class PDFQuestionAnswerer:
```

```
    def __init__(self, claude_api_key):
```

```
...
```

- **PDFQuestionAnswerer**: A class that handles processing PDFs and answering questions based on their contents.

- `__init__`: The constructor method initializes the class, setting up necessary tools (like the Claude API for answering questions and the Hugging Face embeddings for text processing).

### `**4. Extract Text from PDF:**`

```
```python
def extract_text_from_pdf(self, pdf_path):
    try:
        reader = PdfReader(pdf_path)
        text = ""
        for page in reader.pages:
            text += page.extract_text() + "\n"
        return text
    ```
```

- `extract_text_from_pdf`: This method reads a PDF file, extracts text from each page, and combines it into a single string of text.

- `PdfReader(pdf_path)`: Reads the PDF file from the given path.

- `page.extract_text()`: Extracts the text from each page of the PDF.

### `**5. Process Multiple PDFs and Create Vector Store:**`

```
```python
def process_pdfs(self, pdf_directory):
    all_texts = []
    pdf_files = [f for f in os.listdir(pdf_directory) if f.endswith('.pdf')]

    if not pdf_files:
        raise FileNotFoundError(f"No PDF files found in directory: {pdf_directory}")
    ```
```

- `process_pdfs`: This method processes multiple PDF files from a directory and prepares them for answering questions.

- `os.listdir(pdf_directory)`: Lists all files in the specified directory.

- `**f.endswith('.pdf')**: Filters only PDF files.`

```
```python
for pdf_file in pdf_files:
    pdf_path = os.path.join(pdf_directory, pdf_file)
    text = self.extract_text_from_pdf(pdf_path)
    if text:
        all_texts.append(text)
```
```

- This loop reads each PDF file, extracts its text, and adds it to a list of all extracted texts.

```
```python
text_splitter = RecursiveCharacterTextSplitter(
    chunk_size=1000,
    chunk_overlap=200,
    length_function=len
)

chunks = text_splitter.create_documents(all_texts)
```
```

- `**RecursiveCharacterTextSplitter**: Splits the long text into smaller chunks, each with a maximum size of 1000 characters, and allows some overlap between chunks (200 characters).`

- `**create_documents(all_texts)**: Creates smaller text documents from the list of all extracted texts.`

```
```python
self.vector_store = FAISS.from_documents(
    chunks,
    self.embeddings
)
```

```
...
```

- `FAISS.from_documents`: Creates a FAISS vector store using the smaller text chunks, and the text embeddings (numerical representations) generated by the Hugging Face model.

- `self.embeddings`: Converts the text into embeddings so that FAISS can store and search them efficiently.

```
```python
```

```
with open('faiss_vector_store.pkl', 'wb') as f:
```

```
    pickle.dump(self.vector_store, f)
```

```
...
```

- This saves the vector store to a file (`faiss_vector_store.pkl`) using `pickle`, so it can be used later without having to process the PDFs again.

```
### **6. Load the Saved Vector Store:**
```

```
```python
```

```
def load_vector_store(self):
```

```
    try:
```

```
        with open('faiss_vector_store.pkl', 'rb') as f:
```

```
            self.vector_store = pickle.load(f)
```

```
...
```

- `load_vector_store`: This method loads the vector store from the saved file.

- `pickle.load(f)`: Loads the vector store from the pickle file so it can be used again.

```
### **7. Get Relevant Chunks for a Query:**
```

```
```python
```

```
def get_relevant_chunks(self, query, k=3):
```

```
    if not self.vector_store:
```

```
        raise ValueError("No PDFs have been processed yet. Call process_pdfs first.")
```

```
    docs = self.vector_store.similarity_search(query, k=k)
```

```
    return [doc.page_content for doc in docs]
```



...

- `get_relevant_chunks`: This method finds the most relevant chunks of text for a given query.
- `self.vector_store.similarity_search(query, k=k)`: Searches the vector store for the top `k` most relevant chunks for the given query.

### **8. Format the Prompt for Claude (AI):**

```
python
```

```
def format_prompt(self, query, relevant_chunks):
```

```
    context = "\n\n".join(relevant_chunks)
```

```
    prompt = f"""\n\nHuman: Here are some relevant passages from the documents:
```

```
{context}
```

```
Based on the passages above, please answer this question: {query}
```

```
If the answer cannot be fully determined from the provided passages, please say so. Include specific references to the source material where possible.
```

```
Assistant:"""
```

```
    return prompt
```

...

- `format_prompt`: This method formats the query and the relevant text chunks into a prompt that will be sent to Claude (AI).

- It organizes the chunks and the question into a friendly format that Claude can understand.

### **9. Ask a Question Using Claude AI:**

```
python
```

```
def ask_question(self, query):
```

```
    if self.vector_store is None:
```

```
        print("Error: No vector store available. Please process PDFs first.")
```

```
        return "Error: No PDFs processed yet."
```

```
    relevant_chunks = self.get_relevant_chunks(query)
```

```
    prompt = self.format_prompt(query, relevant_chunks)
```

```

response = self.anthropic.messages.create(
    model="claude-3-sonnet-20240229",
    max_tokens=1000,
    temperature=0,
    messages=[{
        "role": "user",
        "content": prompt
    }]
)

```

```

    return response.content[0].text
...

```

- `ask_question`: This method asks Claude (AI) to answer a question based on the relevant chunks of text.
- It first checks if the vector store is available, then formats the prompt and sends it to Claude using the `Anthropic API`.
- `response.content[0].text`: Extracts the answer from the Claude API's response.

### `10. Streamlit User Interface (UI):`

```

python
def main():
    try:
        claude_api_key = os.getenv("CLAUDE_API_KEY")
        pdf_dir = os.getenv("PDF_DIR", "./pdfs")
    ...

```

- `main`: The main entry point for the program. It initializes and runs the Streamlit app.
- It loads the API key (`CLAUDE_API_KEY`) and the PDF directory (`PDF_DIR`).

```

python

```

```
qa_system = PDFQuestionAnswerer(claude_api_key)
qa_system.load_vector_store()
...
```

- Initializes the `PDFQuestionAnswerer` object and loads the existing vector store if available.

```
```python
uploaded_files = st.sidebar.file_uploader("Upload PDF files", type="pdf", accept_multiple_files=True)
...
```

- Creates a file uploader in the Streamlit sidebar that allows users to upload multiple PDF files.

```
```python
if uploaded_files:
    pdf_directory = "./temp_pdfs"
    os.makedirs(pdf_directory, exist_ok=True)
...
```

- If PDF files are uploaded, it creates a temporary directory to store them.

```
```python
for uploaded_file in uploaded_files:
    file_path = os.path.join(pdf_directory, uploaded_file.name)
    with open(file_path, "wb") as f:
        f.write(uploaded_file.read())
...
```

- Saves the uploaded files to the `temp\_pdfs` directory.

```
```python
if st.sidebar.button("Process PDFs") and qa_system.vector_store is None:
    qa_system.process_pdfs(pdf_directory)
    st.sidebar.success("PDFs processed successfully.")
```

```
qa_system.load_vector_store()
'''
```

- When the "Process PDFs" button is clicked, it processes the PDFs and saves the vector store.

```
```python
query = st.text_input("Ask a question about the PDFs:")
if query:
    answer = qa_system.ask_question(query)
    st.write("### Answer:")
    st.write(textwrap.fill(answer, width=80))
'''
```

- Takes a user query and asks Claude (AI) to answer based on the processed PDFs.

---

### \*\*Summary:\*\*

This code:

- Allows users to upload PDF files.
- Extracts text from the PDFs and splits it into chunks.
- Creates a vector store (FAISS index) from the text chunks.
- Saves and loads the vector store to preserve the processed data across interactions.
- Uses Claude (AI) to answer questions based on the content of the PDFs.

Let me know if you need further clarification!

